
Anomaly Detection

Using Python

“Anomalies” Difficult to Characterize

- Define as “Very Different but Rare”?
 - \Rightarrow outlier detection
 - Not applicable to security -- persistent threats
 - How about “New and Not Normal”?
 - \Rightarrow novelty detection
 - How “normal” is your definition of normal?
 - Is the “Anomaly” Worth Reporting?
-

Unique Attributes of the Problem

- Traditional “Detection” Problems:
 - Multi-classification
 - #1: train on labeled data \Rightarrow label new data
 - #2: cluster unlabeled data \Rightarrow label new data
 - Anomaly Detection:
 - One-class classification
 - No known correct model for “different”
 - Detect them anyway as they arise
-

Despite Difficulties, Successful Area


- Fraud Detection, Network Security, Flight Safety, ...
 - We'll Try to Address Previous Points from a Practical Perspective
-

Credit Card Fraud Example

- Scenario:
 - Merchant purchases stolen #'s to buy own goods
 - First-time offender
 - Uses all the stolen cards at once
 - Our Approach:
 - We have historic data of merchant transaction rates
 - We'll detect abnormally **high** transaction rates
-

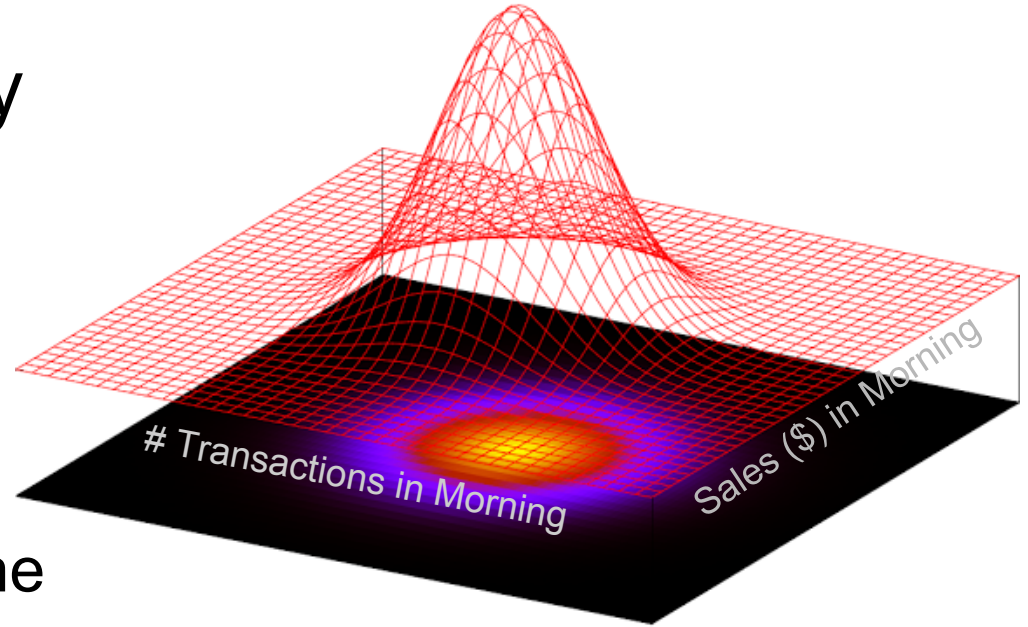
Steps in Python

1. Gather purchase history for the merchant
 - a. # of purchases per day
 - b. # purchases per times of day
2. Approximate purchase data as Gaussian
3. On new day, determine if anomalously high by computing p-value



Two Features Are Easy to Interpret

- Include Total Daily Sales (\$)
- Easy to Visualize
 - 3-D Histograms
 - Colorize in the plane

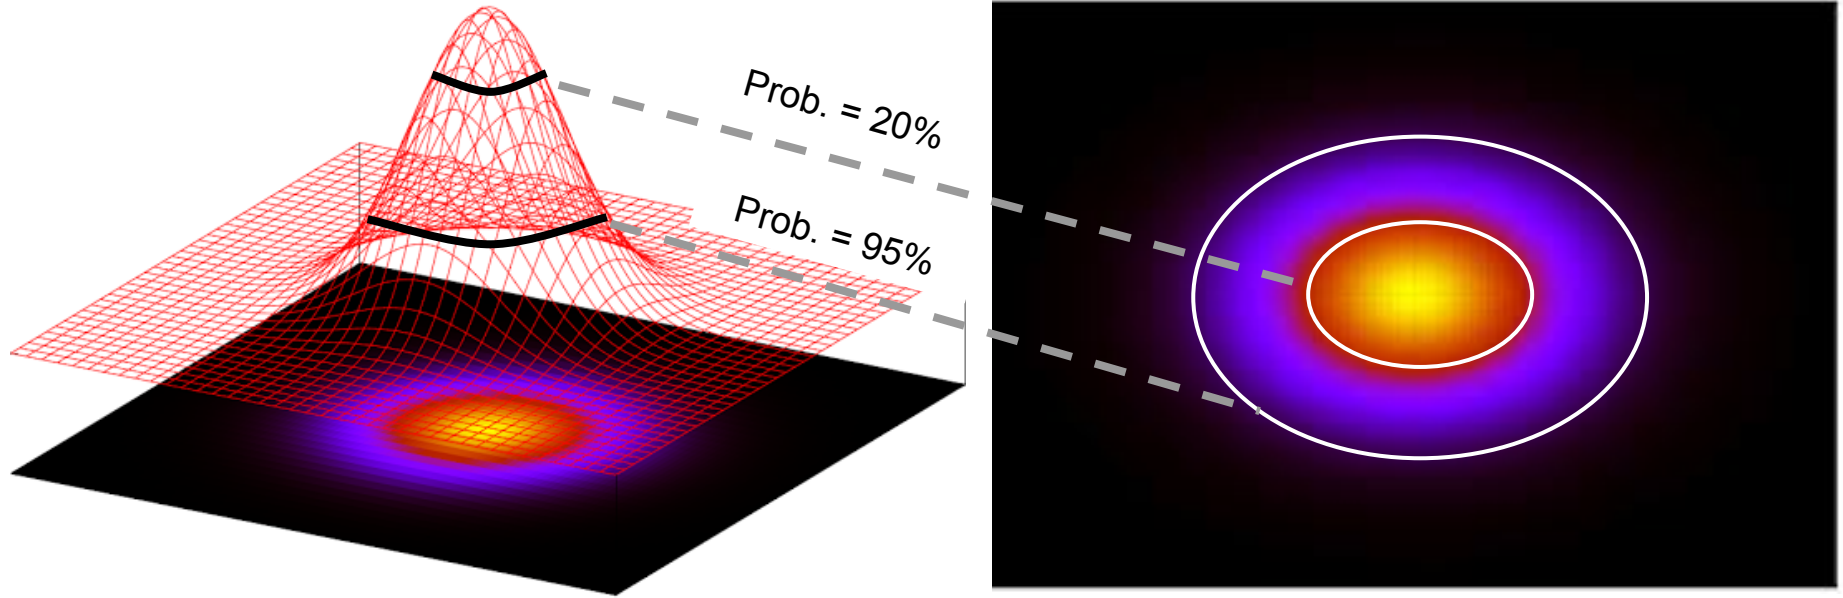


| High Freq

Low Freq |

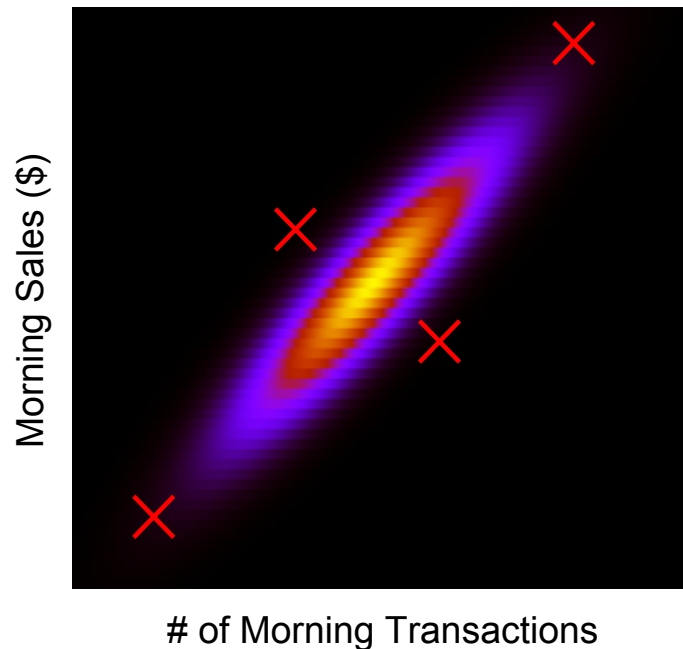


Levels Sets \Rightarrow Probabilities




Which Points Indicate Fraud?

- Same Level Set
 - Level set probability $\sim 99\%$
 - All anomalous
- What Is the Impact of Using Regression?



Robust Distribution Fitting in Python

- Fitting Gaussians in Noisy Data Is Challenging
- Robust: Find Pts that Minimize Area of X% Level Set
- Marks Distant Points as Outliers (Uses Probability)




Credit Card Fraud Example



- Scenario:
 - Two merchants purchase stolen #'s to buy goods
 - We lack historic data for the merchants
 - Have data from many other similar merchants
 - Our Approach:
 - Perform a robust estimation of # customers + sales
 - Detect fraudulent merchants as an outliers
-

Steps in Python

1. Gather sales and transaction data from all merchants
2. Run robust covariance estimation
3. Find those merchants that are outliers



Non-Gaussian Data Is Typical



Mixture Models:


- Probabilistic
- GMMs: sensitive to outliers, need to set K
- Other distributions are robust
- # components can be set by several methods

Distance-Based Methods:

- Non-probabilistic
- Easy-to-understand
- DBSCAN: clusters while searching for outliers

DBSCAN: Density Based Clustering

- Point Classification:
 - Core: satisfies density
 - Boundary: connected to a core point
 - Noise: not connected to a core point
- Key Is Efficient Implementation




Credit Card Fraud Example

- Scenario:
 - Two merchants purchase stolen #'s to buy goods
 - We lack historic data for the merchants
 - Have data from many other similar merchants
 - Our Approach:
 - Apply DBSCAN to # customers + sales
 - Detect fraudulent merchants as an outliers
-

Steps in Python

1. Gather sales and transaction data from all merchants
2. Run DBSCAN
3. Find those merchants that are outliers (noise)



Other Things to Consider...

- Cost of False Positives: What Is the Impact of Accusing a Merchant of Fraud?
 - Timescales: Some Anomalies Are Only Visible over Long Periods
 - Ground Truth: How Can You Take Advantage of Past True or False Positives?
-

Final Remarks

- **Successful Anomaly Detection Requires:**
 - Knowing or robustly estimating normal behavior
 - Knowing which low-probability events to ignore
 - Knowing what part of your data to examine
 - **Many, Many Algorithms Exist**
 - We focused on one probabilistic approach
 - Distance-based heuristics are also used
 - Not covered: 1-class SVMs and many others...
-